

# Exercises on Machine Learning Tools

STSM2634

2025-05-19

```
library(Ecdat)

## Warning: package 'Ecdat' was built under R version 4.3.3
## Loading required package: Ecfun
## Warning: package 'Ecfun' was built under R version 4.3.3
##
## Attaching package: 'Ecfun'
## The following object is masked from 'package:base':
##
##      sign
##
## Attaching package: 'Ecdat'
## The following object is masked from 'package:datasets':
##
##      Orange

library(tidyverse)

## — Attaching core tidyverse packages ————— tidyverse
2.0.0 —
## ✓ dplyr      1.1.2      ✓ readr      2.1.4
## ✓ forcats    1.0.0      ✓ stringr    1.5.0
## ✓ ggplot2    3.4.2      ✓ tibble     3.2.1
## ✓ lubridate  1.9.2      ✓ tidyr      1.3.0
## ✓ purrr      1.0.1

## — Conflicts —————
tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
conflicts to become errors

library(e1071)
library(lattice)
library(AER)

## Warning: package 'AER' was built under R version 4.3.3
```

```
## Loading required package: car
## Warning: package 'car' was built under R version 4.3.1
## Loading required package: carData
##
## Attaching package: 'carData'
##
## The following object is masked from 'package:Ecdat':
##
##     Mroz
##
## Attaching package: 'car'
##
## The following object is masked from 'package:dplyr':
##
##     recode
##
## The following object is masked from 'package:purrr':
##
##     some
##
## Loading required package: lmtest
## Warning: package 'lmtest' was built under R version 4.3.1
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##     as.Date, as.Date.numeric
##
## Loading required package: sandwich
## Loading required package: survival

library(neuralnet)

##
## Attaching package: 'neuralnet'
##
## The following object is masked from 'package:dplyr':
##
##     compute

library(MASS)

##
## Attaching package: 'MASS'
##
```

```
## The following object is masked from 'package:dplyr':
##
##   select
##
## The following object is masked from 'package:Ecdat':
##
##   SP500
```

## Q1. Logistic Regression on Binary Outcome (SwissLabor dataset).

```
data("SwissLabor")
str(SwissLabor)

## 'data.frame':   872 obs. of  7 variables:
## $ participation: Factor w/ 2 levels "no","yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ income      : num  10.8 10.5 11 11.1 11.1 ...
## $ age         : num   3 4.5 4.6 3.1 4.4 4.2 5.1 3.2 3.9 4.3 ...
## $ education   : num   8 8 9 11 12 12 8 8 12 11 ...
## $ youngkids   : num   1 0 0 2 0 0 0 0 0 0 ...
## $ oldkids     : num   1 1 0 0 2 1 0 2 0 2 ...
## $ foreign     : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...

# Fit logistic regression model
modell1 <- glm(participation ~ youngkids + income, data = SwissLabor, family =
binomial)

# Create prediction grid
newdata1 <- data.frame(youngkids = seq(min(SwissLabor$youngkids),
max(SwissLabor$youngkids), length.out = 100),
income = mean(SwissLabor$income))

newdata1$predicted_prob <- predict(modell1, newdata = newdata1, type =
"response")

# Plot
ggplot(SwissLabor, aes(x = youngkids, y = as.numeric(participation) - 1)) +
  geom_jitter(height = 0.05, width = 0.2, alpha = 0.5) +
  geom_line(data = newdata1, aes(x = youngkids, y = predicted_prob), color =
"blue", size = 1.2) +
  labs(title = "Logistic Regression: Labor Participation ~ Fertility",
x = "Fertility", y = "Predicted Probability") +
  theme_minimal()

## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



## Q2. Poisson Regression for Count Data (warpbreaks dataset).

```
data("warpbreaks")
```

```
# Fit Poisson regression model
```

```
model2 <- glm(breaks ~ wool + tension, data = warpbreaks, family = poisson)
```

```
# Predict fitted values
```

```
warpbreaks$fit <- predict(model2, type = "response")
```

```
# Plot
```

```
ggplot(warpbreaks, aes(x = tension, y = breaks, color = wool)) +  
  geom_point() +  
  geom_line(aes(y = fit, group = wool), size = 1.2) +  
  labs(title = "Poisson GLM: Breaks ~ Wool + Tension",  
       x = "Tension", y = "Observed & Fitted Breaks") +  
  theme_minimal()
```



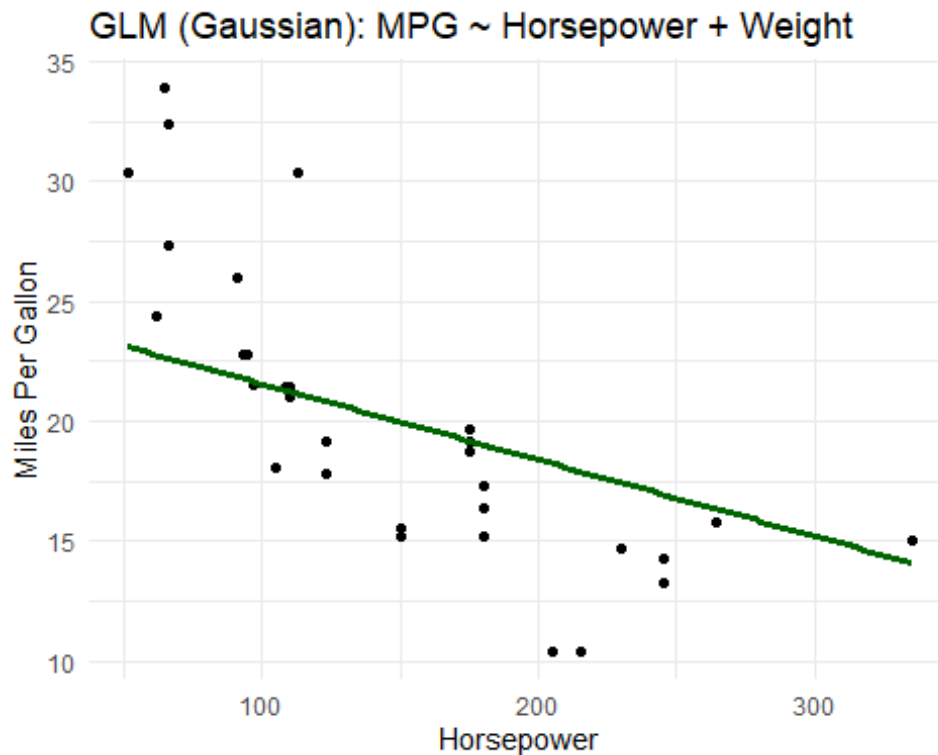
### Q3. Ordinary Linear Regression (mtcars dataset).

```
data("mtcars")
```

```
# Fit GLM (default family = Gaussian)
model3 <- glm(mpg ~ hp + wt, data = mtcars)
```

```
# Predict over horsepower range at fixed weight
newdata3 <- data.frame(hp = seq(min(mtcars$hp), max(mtcars$hp), length.out =
100),
                        wt = mean(mtcars$wt))
newdata3$predicted <- predict(model3, newdata = newdata3)
```

```
# Plot
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point() +
  geom_line(data = newdata3, aes(x = hp, y = predicted), color = "darkgreen",
size = 1.2) +
  labs(title = "GLM (Gaussian): MPG ~ Horsepower + Weight",
       x = "Horsepower", y = "Miles Per Gallon") +
  theme_minimal()
```



Q4. Applying SVM on Wages data (Binary classification) to model the income status based on the education, experience, marital status, and sex.

```
data(Wages)
str(Wages)

## 'data.frame':  4165 obs. of  12 variables:
## $ exp      : int  3 4 5 6 7 8 9 30 31 32 ...
## $ wks      : int  32 43 40 39 42 35 32 34 27 33 ...
## $ bluecol  : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 2 2 2 ...
## $ ind      : int  0 0 0 0 1 1 1 0 0 1 ...
## $ south    : Factor w/ 2 levels "no","yes": 2 2 2 2 2 2 2 1 1 1 ...
## $ smsa     : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ married  : Factor w/ 2 levels "no","yes": 2 2 2 2 2 2 2 2 2 2 ...
## $ sex      : Factor w/ 2 levels "female","male": 2 2 2 2 2 2 2 2 2 2 ...
## $ union    : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 2 ...
## $ ed       : int  9 9 9 9 9 9 9 11 11 11 ...
## $ black    : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ lwage    : num  5.56 5.72 6 6 6.06 ...

any(is.na(Wages))

## [1] FALSE

# Create binary target: high wage
mean_wage <- mean(Wages$lwage, na.rm = TRUE)
```

```

Wages_clean <- na.omit(Wages) # Remove missing values, if any
Wages_clean$high_wage <- as.factor(Wages_clean$lwage > mean_wage)

# Split into training and test sets (e.g., 70/30)
set.seed(123)
n <- nrow(Wages_clean)
train_index <- sample(1:n, size = 0.7 * n)
train_data <- Wages_clean[train_index, ]
test_data <- Wages_clean[-train_index, ]

# Fit SVM
model2 <- svm(high_wage ~ exp + ed + married + sex, data = train_data, kernel
= "radial", probability = TRUE)

# Training accuracy
pred2 <- predict(model2, test_data)
acc2 <- mean(pred2 == test_data$high_wage)
cat("Training Accuracy (Wages):", round(acc2, 4), "\n")

## Training Accuracy (Wages): 0.7152

# Dummy input
dummy2 <- data.frame(exp = 10, ed = 14, married = factor("yes", levels =
levels(Wages_clean$married)), sex = factor("male", levels =
levels(Wages_clean$sex)) )
pred_dummy2 <- predict(model2, dummy2, probability = TRUE)
cat("Dummy Prediction (Wages):", as.character(pred_dummy2), "\n")

## Dummy Prediction (Wages): FALSE

```

**Q4. Applying SVM on SwissLabor data (Binary classification) to model the participation status based on the education, age, and income.**

```

# Load data
data(SwissLabor)
str(SwissLabor)

## 'data.frame':    872 obs. of  7 variables:
## $ participation: Factor w/ 2 levels "no","yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ income      : num  10.8 10.5 11 11.1 11.1 ...
## $ age         : num   3 4.5 4.6 3.1 4.4 4.2 5.1 3.2 3.9 4.3 ...
## $ education   : num   8 8 9 11 12 12 8 8 12 11 ...
## $ youngkids   : num   1 0 0 2 0 0 0 0 0 0 ...
## $ oldkids     : num   1 1 0 0 2 1 0 2 0 2 ...
## $ foreign     : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...

Swiss_clean <- na.omit(SwissLabor) # Remove missing values, if any

# Target: participate (already binary: yes/no)
Swiss_clean$participate <- factor(Swiss_clean$participation)

```

```

# Split into train/test
set.seed(123)
n <- nrow(Swiss_clean)
train_idx <- sample(1:n, size = 0.7 * n)
train_data <- Swiss_clean[train_idx, ]
test_data <- Swiss_clean[-train_idx, ]

# Fit SVM
model_swiss <- svm(participation ~ age + income + education,
                   data = train_data, kernel = "radial", probability = TRUE)

# Test accuracy
pred_swiss <- predict(model_swiss, test_data)
acc_swiss <- mean(pred_swiss == test_data$participate)
cat("Test Accuracy (SwissLabor):", round(acc_swiss, 4), "\n")

## Test Accuracy (SwissLabor): 0.6221

# Dummy input
dummy_swiss <- data.frame(age = 35, income = 5, education = 12)
pred_dummy_swiss <- predict(model_swiss, dummy_swiss, probability = TRUE)
cat("Dummy Prediction (SwissLabor):", as.character(pred_dummy_swiss), "\n")

## Dummy Prediction (SwissLabor): no

```

## Q5. Housing Price Categorization using Boston data from MASS package with SVM.

```

data(Boston)
str(Boston)

## 'data.frame':    506 obs. of  14 variables:
## $ crim      : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn        : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus     : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas      : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox       : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524
0.524 ...
## $ rm        : num  6.58 6.42 7.18 7 7.15 ...
## $ age       : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis       : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad       : int   1 2 2 3 3 3 5 5 5 5 ...
## $ tax       : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio   : num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black     : num  397 397 393 395 397 ...
## $ lstat     : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv      : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...

```



```

# Create binary target
median_price <- median(Boston$medv)
Boston$expensive <- as.factor(Boston$medv > median_price)

# Split into train/test
set.seed(123)
n <- nrow(Boston)
train_idx <- sample(1:n, size = 0.7 * n)
train_data <- Boston[train_idx, ]
test_data <- Boston[-train_idx, ]

# Fit SVM
model_boston <- svm(expensive ~ lstat + rm + crim, data = train_data, kernel
= "radial", probability = TRUE)

# Test accuracy
pred_boston <- predict(model_boston, test_data)
acc_boston <- mean(pred_boston == test_data$expensive)
cat("Test Accuracy (Boston):", round(acc_boston, 4), "\n")

## Test Accuracy (Boston): 0.8553

# Dummy input
dummy_boston <- data.frame(lstat = 5, rm = 7, crim = 0.05)
pred_dummy_boston <- predict(model_boston, dummy_boston, probability = TRUE)
cat("Dummy Prediction (Boston):", as.character(pred_dummy_boston), "\n")

## Dummy Prediction (Boston): TRUE

```

Q6. Applying ANN on SwissLabor data (Binary classification) to model the participation status based on the education, age, and income.

```

data(SwissLabor)
str(SwissLabor)

## 'data.frame':    872 obs. of  7 variables:
## $ participation: Factor w/ 2 levels "no","yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ income       : num  10.8 10.5 11 11.1 11.1 ...
## $ age          : num   3 4.5 4.6 3.1 4.4 4.2 5.1 3.2 3.9 4.3 ...
## $ education    : num   8 8 9 11 12 12 8 8 12 11 ...
## $ youngkids    : num   1 0 0 2 0 0 0 0 0 0 ...
## $ oldkids      : num   1 1 0 0 2 1 0 2 0 2 ...
## $ foreign      : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...

Swiss_clean <- na.omit(SwissLabor)
Swiss_clean$participate_num <- ifelse(Swiss_clean$participation == "yes", 1,
0)

# Normalize inputs
Swiss_clean$income <- scale(Swiss_clean$income)

```

```

Swiss_clean$education <- scale(Swiss_clean$education)
Swiss_clean$age <- scale(Swiss_clean$age)

# Train-test split
set.seed(123)
n <- nrow(Swiss_clean)
train_idx <- sample(1:n, size = 0.7 * n)
train_data <- Swiss_clean[train_idx, ]
test_data <- Swiss_clean[-train_idx, ]

# Train neural network
modell1 <- neuralnet(participate_num ~ age + income + education,
                     data = train_data, hidden = c(3), linear.output = FALSE)

plot(modell1)
# Predict on test set
test_pred1 <- compute(modell1, test_data[, c("age", "income",
"education")])$net.result
test_class1 <- ifelse(test_pred1 > 0.5, 1, 0)
acc1 <- mean(test_class1 == test_data$participate_num)
cat("Test Accuracy (SwissLabor):", round(acc1, 4), "\n")

## Test Accuracy (SwissLabor): 0.6374

# Dummy input
dummy1 <- data.frame(
  age = scale(35, attr(Swiss_clean$age, "scaled:center"),
attr(Swiss_clean$age, "scaled:scale")),
  income = scale(5, attr(Swiss_clean$income, "scaled:center"),
attr(Swiss_clean$income, "scaled:scale")),
  education = scale(12, attr(Swiss_clean$education, "scaled:center"),
attr(Swiss_clean$education, "scaled:scale"))
)

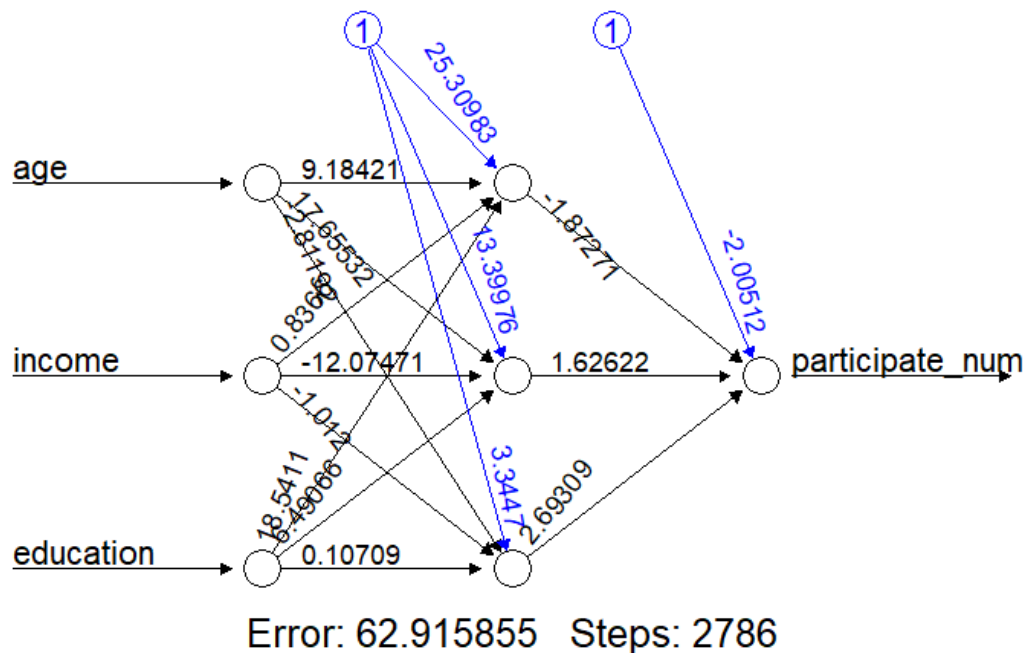
pred_dummy1 <- compute(modell1, dummy1)$net.result
cat("Dummy Prediction (SwissLabor):", round(pred_dummy1, 4), "\n")

## Dummy Prediction (SwissLabor): 0.0952

# Convert probability to class label
class_dummy1 <- ifelse(pred_dummy1 > 0.5, "yes", "no")
cat("Dummy Prediction (SwissLabor):", class_dummy1, "\n")

## Dummy Prediction (SwissLabor): no

```



Q7. Fit an ANN model on the mtcars dataset to predict the mpg values based on the hp, wt, disp values.

```
data(mtcars)
str(mtcars)

## 'data.frame':  32 obs. of  11 variables:
## $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num   6  6  4  6  8  6  8  4  4  6 ...
## $ disp: num  160 160 108 258 360 ...
## $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num   16.5 17 18.6 19.4 17 ...
## $ vs  : num    0  0  1  1  0  1  0  1  1  1 ...
## $ am  : num    1  1  1  0  0  0  0  0  0  0 ...
## $ gear: num    4  4  4  3  3  3  3  4  4  4 ...
## $ carb: num    4  4  1  1  2  1  4  2  2  4 ...

# Normalize all numeric columns
mtcars_scaled <- as.data.frame(scale(mtcars))

# Train-test split
set.seed(123)
n <- nrow(mtcars_scaled)
```

```

train_idx <- sample(1:n, size = 0.7 * n)
train_data <- mtcars_scaled[train_idx, ]
test_data <- mtcars_scaled[-train_idx, ]

# Train model
model2 <- neuralnet(mpg ~ hp + wt + disp,
                    data = train_data, hidden = 5, linear.output = TRUE)

plot(model2)
# Predict on test set
test_pred2 <- compute(model2, test_data[, c("hp", "wt", "disp")])$net.result
rmse2 <- sqrt(mean((test_pred2 - test_data$mpg)^2))
cat("Test RMSE (mtcars):", round(rmse2, 4), "\n")

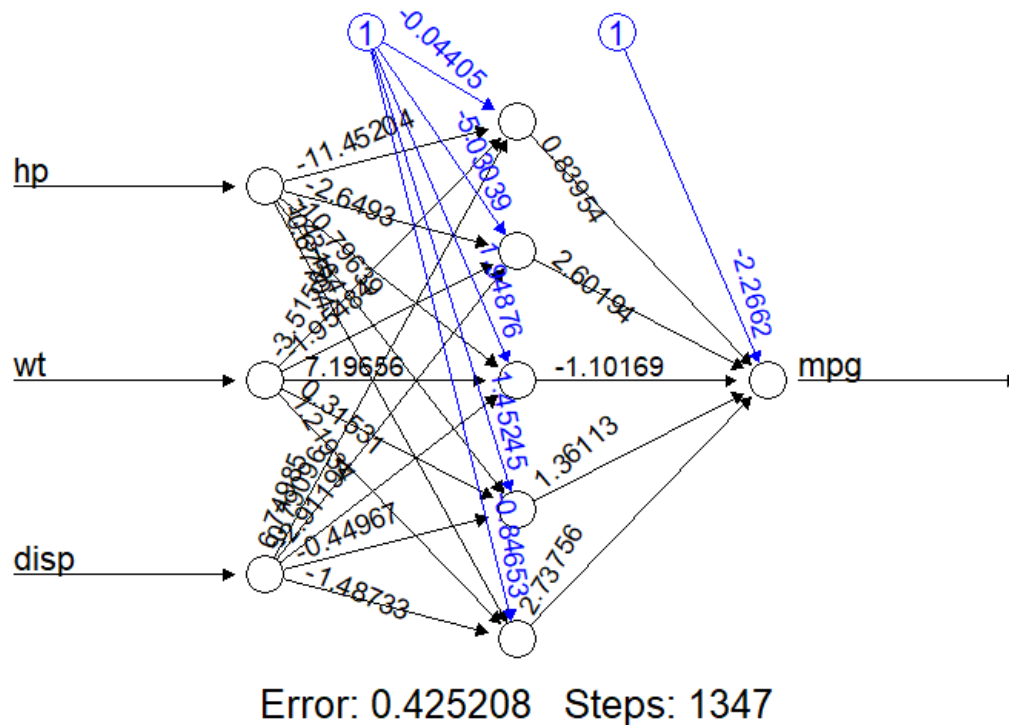
## Test RMSE (mtcars): 0.4037

# Dummy input
dummy2 <- data.frame(
  hp = (120 - mean(mtcars$hp)) / sd(mtcars$hp),
  wt = (2.8 - mean(mtcars$wt)) / sd(mtcars$wt),
  disp = (180 - mean(mtcars$disp)) / sd(mtcars$disp)
)

pred_dummy2 <- compute(model2, dummy2)$net.result
# Rescale prediction
pred_mpg <- pred_dummy2 * sd(mtcars$mpg) + mean(mtcars$mpg)
cat("Dummy Prediction (mpg):", round(pred_mpg, 2), "\n")

## Dummy Prediction (mpg): 17.76

```



Q8. Analyze relationships between wage-related numeric variables with PCA.

```
# Load and clean data
data(Wages)
Wages_clean <- na.omit(Wages)

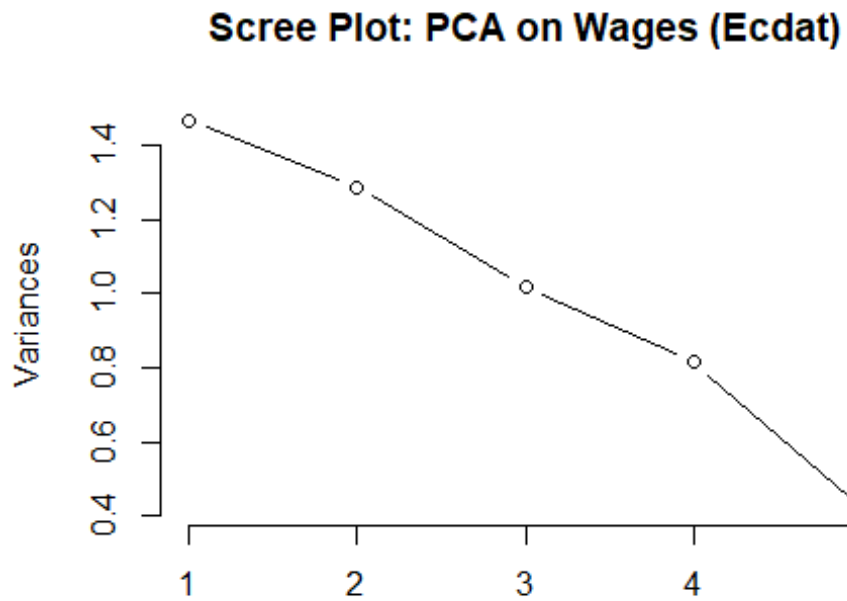
# Select numeric variables for PCA
vars <- Wages_clean[, c("exp", "wks", "ind", "ed", "lwage")]

# Perform PCA
pca_wages <- prcomp(vars, center = TRUE, scale. = TRUE)

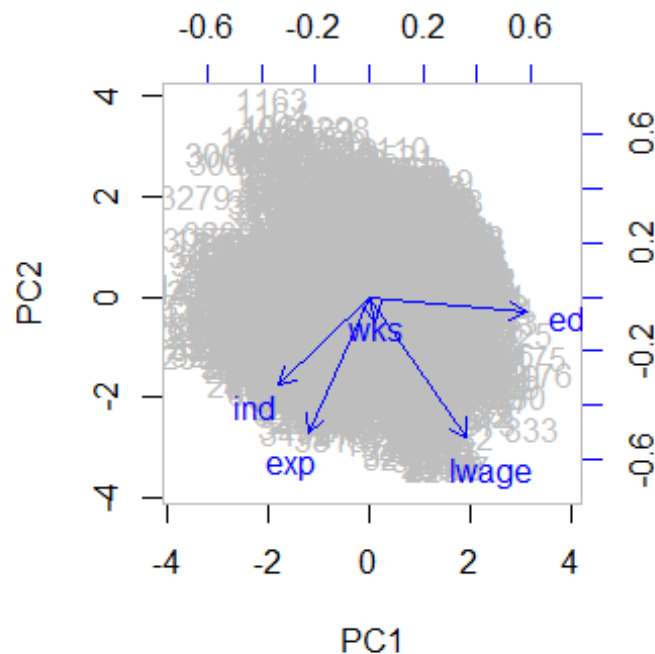
# View variance explained
summary(pca_wages)

## Importance of components:
##              PC1      PC2      PC3      PC4      PC5
## Standard deviation  1.2105 1.1345 1.0076 0.9025 0.6465
## Proportion of Variance 0.2931 0.2574 0.2030 0.1629 0.0836
## Cumulative Proportion 0.2931 0.5505 0.7535 0.9164 1.0000
```

```
# Scree plot (plot the principal components according to the variances explained by them)  
plot(pca_wages, type = "l", main = "Scree Plot: PCA on Wages (Ecdat)")
```



```
# Biplot for visualizing variable loadings and scores  
biplot(pca_wages, scale = 0, col = c("gray", "blue"))
```



## Q9. Reduce socio-economic predictors of labor participation using PCA.

```
data(SwissLabor)
str(SwissLabor)

## 'data.frame':  872 obs. of  7 variables:
## $ participation: Factor w/ 2 levels "no","yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ income      : num  10.8 10.5 11 11.1 11.1 ...
## $ age         : num   3 4.5 4.6 3.1 4.4 4.2 5.1 3.2 3.9 4.3 ...
## $ education   : num   8 8 9 11 12 12 8 8 12 11 ...
## $ youngkids   : num   1 0 0 2 0 0 0 0 0 0 ...
## $ oldkids     : num   1 1 0 0 2 1 0 2 0 2 ...
## $ foreign     : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...

Swiss_clean <- na.omit(SwissLabor)

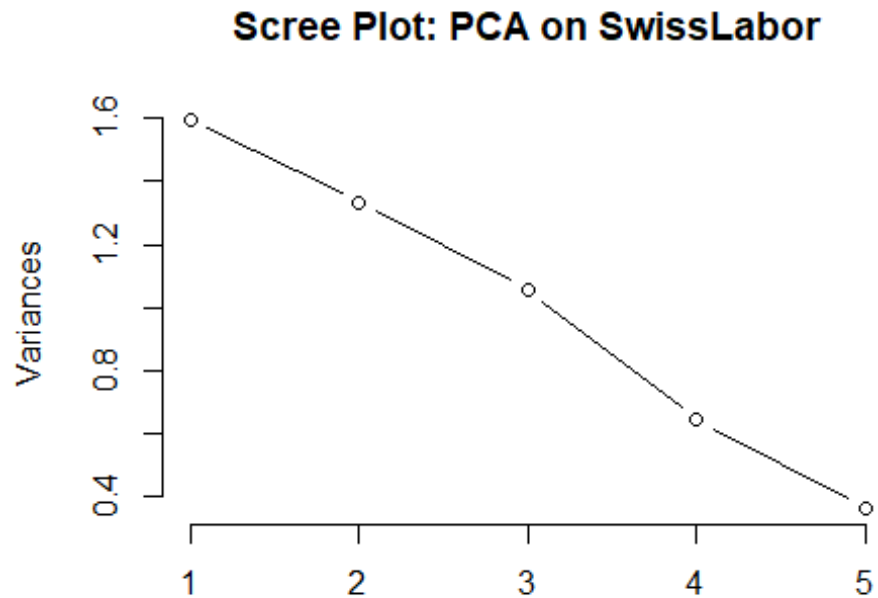
# Use relevant numeric predictors
vars <- SwissLabor[, c("income", "age", "education", "youngkids", "oldkids")]

# Perform PCA
pca_swiss <- prcomp(vars, center = TRUE, scale. = TRUE)

# Summary of variance explained
summary(pca_swiss)
```

```
## Importance of components:
##
##          PC1    PC2    PC3    PC4    PC5
## Standard deviation  1.2633 1.1554 1.0289 0.8032 0.60423
## Proportion of Variance 0.3192 0.2670 0.2117 0.1290 0.07302
## Cumulative Proportion 0.3192 0.5862 0.7979 0.9270 1.00000

# Scree plot
plot(pca_swiss, type = "l", main = "Scree Plot: PCA on SwissLabor")
```



```
# Biplot: PC1 vs PC2
biplot(pca_swiss, scale = 0, col = c("gray40", "blue"), cex = 0.5)
```



